

Extra Review



UNIVERSITY OF
SAN FRANCISCO

Abbie M Popa

BSDS 100 - Intro to Data Science with R



- Something Old
 - Tidy Data
 - Suggestions for extra practice (functions, plotting, string analysis)
- Something New
 - Web-scraping
 - Machine Learning
- A wrap-up activity
- Course review

Request for Extra Review



- Two Approaches:
 - 1 Formal Definition
 - 2 Practical Definition

Tidy Data - The Shelter Animal Example



species	wgt_admitted	wgt_adopted	supp
"cat "	8	15	"turkey"
"cat "	9	11	"kibbles"
"dog"	18	27	"turkey"

These data are not tidy, let's find out why!



Jeff Leek, *The Elements of Data Analytic Style*

- 1 Each variable should be one column.
- 2 Each different observation should be a different row.
- 3 There should be one table for each "kind" of variable.
- 4 If you have multiple tables, they should include a column that allows them to be linked



species	wgt_admitted	wgt_adopted	supp
"cat "	8	15	"turkey"
"cat "	9	11	"kibbles"
"dog"	18	27	"turkey"

❶ Is each variable one column?



species	wgt_admitted	wgt_adopted	supp
"cat "	8	15	"turkey"
"cat "	9	11	"kibbles"
"dog"	18	27	"turkey"

- 1 Is each variable one column? **No! Weight is split across two columns, so is time of measurement**
- 2 Is each observation one row?



species	wgt_admitted	wgt_adopted	supp
"cat "	8	15	"turkey"
"cat "	9	11	"kibbles"
"dog"	18	27	"turkey"

- 1 Is each variable one column? No! Weight is split across two columns, so is time of measurement
- 2 Is each observation one row? No! Each row contains two observations of the animals weight
- 3 Is there one table for each type of variable?



species	wgt_admitted	wgt_adopted	supp
"cat "	8	15	"turkey"
"cat "	9	11	"kibbles"
"dog"	18	27	"turkey"

- ❶ Is each variable one column? No! Weight is split across two columns, so is time of measurement
- ❷ Is each observation one row? No! Each row contains two observations of the animals weight
- ❸ Is there one table for each type of variable? **n/a, only one table**
- ❹ If there are multiple tables, is there a linking column?

Tidy Animals - Formal Definition Approach



species	wgt_admitted	wgt_adopted	supp
"cat "	8	15	"turkey"
"cat "	9	11	"kibbles"
"dog"	18	27	"turkey"

- ❶ Is each variable one column? No! Weight is split across two columns, so is time of measurement
- ❷ Is each observation one row? No! Each row contains two observations of the animals weight
- ❸ Is there one table for each type of variable? n/a, only one table
- ❹ If there are multiple tables, is there a linking column? n/a, only one table



- There are at least three approaches to how to decide how many tables to split data across
- This will be covered in more detail in the databases class, CS 333



Format the data so that you can complete the goal you have in mind

- I want to make a plot showing what influences change in weight over time
- I want to formally test what affects weight using a linear model



- Functions:

- <https://r4ds.had.co.nz/index.html> chapter 19
- <https://exercism.io/tracks/r> - requires login, but free
- <https://www.hackerrank.com> - requires login, but free, select R from dropdown

- Plotting:

- <https://r4ds.had.co.nz/index.html> chapter 3
- <http://www.cookbook-r.com/Graphs/> type their code, try on your own data

- String Analysis:

- <https://r4ds.had.co.nz/index.html> chapter 14
- <https://regexr.com/>, just remember, R requires the double backslash where most languages use the single backslash!



But remember, the best way to practice R (or any programming language) is to complete projects!

- Enter a kaggle competition <https://www.kaggle.com/>
- Enter a driven data competition <https://www.drivendata.org/>
- Solve your own problem with data! Then share it...

<https://github.com/>

<https://medium.com/>

Your personal website!

- Examples:
 - <https://pudding.cool/>
 - <https://flowingdata.com/>
 - <https://towardsdatascience.com/>

New Topics (Very Brief)



- Types:
 - Some websites have an API, where you can directly request data be served to you in a convenient format (e.g., CSV)
 - Most websites are written in `html`, a structured way to format and display text
 - Some website rely heavily on `javascript`, which is code that executes in your browser
- We will focus on the second, which is common and doable



- Examine the web page using the browser's developer tools
- Import and parse the webpage using the `rvest` package
- Example, build a corpus of legal terms to see how many lawyers are using our web app



- Tutorial

<https://www.analyticsvidhya.com/blog/2017/03/beginners-guide-on-web-scraping-in-r-using-rvest-with-hands-on-knowledge/>

- Major sticking points:

- Javascript (requires Selenium)
- If they change their webpage, you'll need to update!

- Data Mining Class: CS 451



- Goal: Use a collected data set, to build a model, that can be used to make predictions about new data
- Many approaches all with different pros and cons (no free lunch!)
- A linear regression example



- Machine Learning is a huge topic!
- Several Types
 - Supervised Learning - data has labels
 - Unsupervised Learning - data does not have labels
 - Regression - Predicted Value is continuous
 - Classification - Predicted Value is categorical
- Different levels of flexibility
 - Linear regression is not very flexible, the output model is always a straight line
 - Other options, such as decision trees, are more flexible



- Workflow

- ① Split data into train and test sets
- ② Train data on the training set
- ③ Check accuracy on test set
- ④ Try with different algorithms or hyperparameters

- Considerations

- More flexible models have a higher risk of "overfitting", that is, they will work well on the training data but not new data
- This is why with more flexible models it is important to have a larger training dataset



- Popular Algorithms

- Linear Regression (seen earlier), not very flexible, but works well with small, normal, data
- Random Forests, get a lot of flexibility but use some tricks to avoid overfitting
- Artificial Neural Networks, self-engineer features, often produce state of the art results, but you need more data



- Free online: <https://www-bcf.usc.edu/~gareth/ISL/>
- Not in \mathbb{R} , but beloved
<https://www.coursera.org/learn/machine-learning>
- Class at USF: MATH 373

Wrap-Up



- In groups, complete the DS in the Wild Song Lyrics activity from github
- We will vote on the quality of their report together



You rate your Uber driver, why not rate your classes?
Thank you!